

KERNEL METHODS



FEATURE MAPS

Example. Non-linear classifiers.


$$x = (x_1, x_2)$$

$$\phi(x) = (1, \sqrt{2}x_1, \sqrt{2}x_2, \sqrt{2}x_1x_2, x_1^2, x_2^2)$$

$$h(x; \theta, \theta_0)$$

$$= \text{sign}(\theta \cdot \phi(x))$$



$$= \text{sign}(\theta_1 + \theta_2\sqrt{2}x_1 + \theta_3\sqrt{2}x_2 + \theta_4\sqrt{2}x_1x_2 + \theta_5x_1^2 + \theta_6x_2^2)$$



FEATURE MAPS

Example. Non-linear classifiers.

$$\begin{aligned}h(x; \theta, \theta_0) \\&= \text{sign}((x_1 - 1)^2 + (x_2 - 2)^2 - 9) \\&= \text{sign}(-4 - 2x_1 - 4x_2 + x_1^2 + x_2^2)\end{aligned}$$



CHALLENGES


High-Dimensional Features.

$$x = (x_1, x_2, \dots, x_{1000}) \in \mathbb{R}^{1000}$$

$$\phi(x) = (1, \dots, \sqrt{2}x_i, \dots, \sqrt{2}x_i x_j, \dots, x_i^2, \dots) \in \mathbb{R}^{501501}$$

Inner Products.

Computing $\phi(x) \cdot \phi(x')$ for $x, x' \in \mathbb{R}^{1000}$ requires about 2,004,000 floating point operations.




KERNEL FUNCTIONS

Fortunately, many inner products simplify nicely.

$$\begin{aligned} K(x, x') &= \phi(x) \cdot \phi(x') \\ &= 1 + 2 \sum_i x_i x'_i + 2 \sum_{i < j} x_i x_j x'_i x'_j + \sum_i x_i^2 x_i'^2 \\ &= 1 + 2(\sum_i x_i x'_i) + (\sum_i x_i x'_i)^2 \\ &= (x \cdot x' + 1)^2 \end{aligned}$$

For $x, x' \in \mathbb{R}^{1000}$, computing this requires only about 2000 floating point operations, less than the 501,501 operations needed for $\phi(x)$.



KERNEL FUNCTIONS

'Infinite dimensional'
positive definite
matrices

Definition.

A function $K: \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$ is a *kernel function* if


1. $K(x, y) = K(y, x)$ for all $x, y \in \mathbb{R}^d$,
2. given $n \in \mathbb{N}$ and $x^{(1)}, x^{(2)}, \dots, x^{(n)} \in \mathbb{R}^d$,
the *Gram matrix* K with entries

$$K_{ij} = K(x^{(i)}, x^{(j)})$$

is positive semidefinite.

Example. $K(x, x') = \phi(x) \cdot \phi(x')$

Can be shown that all kernel
functions are of this form!



EXAMPLES

Linear Kernel.

$$K(x, x') = x \cdot x'$$


Polynomial Kernel.


$$K(x, x') = (x \cdot x' + 1)^k$$

Radial Basis Kernel.


$$K(x, x') = \exp\left(-\frac{1}{2}\|x - x'\|^2\right)$$

Feature map $\phi(x)$ is
infinite dimensional!






KERNEL TRICK



KERNEL TRICK

The **kernel trick** refers to the strategy of converting a learning algorithm and the resulting predictor into ones that involve only the computation of the **kernel** $K(x, x') = \phi(x) \cdot \phi(x')$ but not of the **feature map** $\phi(x)$.




SUPPORT VECTOR MACHINES

Learning.

$$\begin{aligned} &\text{maximize } \sum_{(x,y)} \alpha_{x,y} - \frac{1}{2} \sum_{(x,y)} \sum_{(x',y')} \alpha_{x,y} \alpha_{x',y'} y y' (x \cdot x') \\ &\text{subject to } \alpha_{x,y} \geq 0 \text{ for all } (x,y) \end{aligned}$$

Prediction.

$$h(x; \theta) = \text{sign}(\theta \cdot x) = \text{sign} \left(\sum_{(x',y')} \alpha_{x',y'} y' (x \cdot x') \right)$$



KERNEL SUPPORT VECTOR MACHINES


Learning.

$$\begin{aligned} &\text{maximize } \sum_{(x,y)} \alpha_{x,y} - \frac{1}{2} \sum_{(x,y)} \sum_{(x',y')} \alpha_{x,y} \alpha_{x',y'} y y' K(x, x') \\ &\text{subject to } \alpha_{x,y} \geq 0 \text{ for all } (x, y) \end{aligned}$$

Prediction.

$$h(x; \theta) = \text{sign}(\theta \cdot \phi(x)) = \text{sign} \left(\sum_{(x',y')} \alpha_{x',y'} y' K(x, x') \right)$$

We may use the linear, polynomial, or radial basis kernels to get different kinds of decision boundaries.



PERCEPTRON

Learning.


1. Initialize $\theta = 0, \theta_0 = 0$.
2. Repeat until no mistakes are found:
Select data $(x, y) \in \mathcal{S}_n$ in sequence:
If $y(\theta^\top x + \theta_0) \leq 0$, then $\theta \leftarrow \theta + yx, \theta_0 \leftarrow \theta_0 + y$.

Prediction.

$$h(x; \theta, \theta_0) = \text{sign}(\theta \cdot x + \theta_0)$$

From the learning algorithm, we see that

$$\theta = \sum_{x,y} \alpha_{x,y} yx \quad \text{and} \quad \theta_0 = \sum_{x,y} \alpha_{x,y} y \quad \text{for some } \alpha_{x,y} \in \mathbb{N}.$$




PERCEPTRON

Learning.

1. Initialize $\theta = 0, \theta_0 = 0$.
2. Repeat until no mistakes are found:
Select data $(x, y) \in \mathcal{S}_n$ in sequence:
If $\sum_{x', y'} \alpha_{x', y'} y y' (\mathbf{x} \cdot \mathbf{x}' + 1) \leq 0$, then $\alpha_{x, y} \leftarrow \alpha_{x, y} + 1$.

Prediction.

$$h(x; \theta, \theta_0) = \text{sign} \left(\sum_{x', y'} \alpha_{x', y'} y' (\mathbf{x} \cdot \mathbf{x}' + 1) \right)$$




KERNEL PERCEPTRON

Learning.


1. Initialize $\alpha = 0$.
2. Repeat until no mistakes are found:
Select data $(x, y) \in \mathcal{S}_n$ in sequence:
If $\sum_{x', y'} \alpha_{x', y'} y y' (K(x, x') + 1) \leq 0$, then $\alpha_{x, y} \leftarrow \alpha_{x, y} + 1$.

Prediction.


$$h(x; \theta, \theta_0) = \text{sign} \left(\sum_{x', y'} \alpha_{x', y'} y' (K(x, x') + 1) \right)$$



KERNEL PERCEPTRON



Data is always separable when using radial basis kernels!



LINEAR REGRESSION

Learning.

Let $n\lambda\alpha_{x,y} = y - \theta \cdot x$ so we have $n\lambda\alpha = Y - X\theta$.

Recall that the exact solution θ satisfies

$$(n\lambda I + X^\top X)\theta = X^\top Y$$

so we may derive the following:


$$X(n\lambda I + X^\top X)\theta = XX^\top Y$$

$$n\lambda(Y - n\lambda\alpha) + XX^\top(Y - n\lambda\alpha) = XX^\top Y$$

$$n\lambda Y - (n\lambda I + XX^\top)n\lambda\alpha = 0$$

$$\alpha = (n\lambda I + K)^{-1}Y$$

Gram matrix K
with entries
 $K_{ij} = x^{(i)} \cdot x^{(j)}$



LINEAR REGRESSION


Prediction.

Moreover,

$$\begin{aligned} n\lambda X^\top \alpha &= X^\top Y - X^\top X \theta \\ &= (n\lambda I + X^\top X) \theta - X^\top X \theta = n\lambda \theta \end{aligned}$$

So $\theta = X^\top \alpha = \sum_{(x', y')} \alpha_{x', y'} x'$. Therefore,

$$y = \theta \cdot x = \sum_{(x', y')} \alpha_{x', y'} x \cdot x'$$



KERNEL LINEAR REGRESSION


Learning.

$$\alpha = (n\lambda I + K)^{-1}Y$$






Gram matrix K
with entries
 $K_{ij} = K(x^{(i)}, x^{(j)})$

Prediction.

$$y = \sum_{(x', y')} \alpha_{x', y'} K(x, x')$$



KERNEL LINEAR REGRESSION



SUMMARY


- Kernel Functions
 - Feature Maps
 - Inner Products
 - Polynomial Kernel
 - Radial Basis Kernel
- Kernel Trick
 - Support Vector Machines
 - Perceptron
 - Linear Regression



CONVOLUTIONAL NEURAL NETWORKS



IMAGE RECOGNITION



Too many parameters.

- $(\text{MNIST } 28 \times 28 \text{ pixels}) * (100 \text{ features}) = 78,400 \text{ params}$

Translation invariance.

- Statistics of a patch (usually) does not depend on its location.






IMAGE RECOGNITION



Solution.

- Learn good features from small image patches.
- Use features to compute better representation of image.
- Down-sample image to reduce dimensionality.
- Repeat.




CONVOLUTION

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature



FEATURE MAP




Image




Feature Map

-1	-1	-1
2	2	2
-1	-1	-1


Filter



POOLING



Convolved
feature


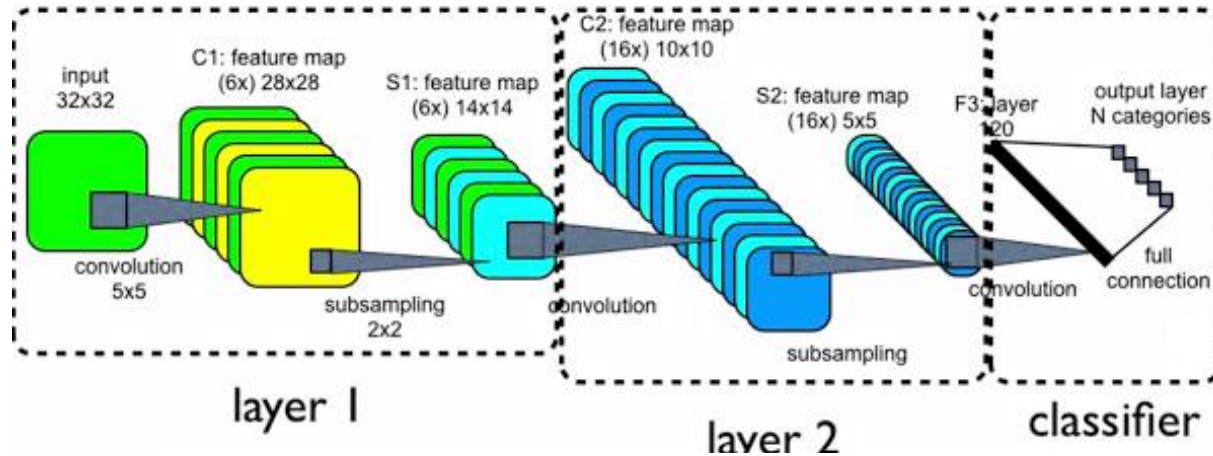


Pooled
feature


Examples. Mean Pooling, Max Pooling.




CONVOLUTIONAL PYRAMID




SHARED WEIGHTS



**Fully-Connected
(Dense)**



**Locally-Connected
with Shared Weights**



DISCUSSION

Question 1.

How is backpropagation and fine-tuning performed?


Question 2.

How to introduce rotation, skew, scale invariance?





TENSORFLOW



ONLINE TUTORIAL

Basic Usage.

- https://www.tensorflow.org/versions/r0.11/get_started/basic_usage.html

Simple MNIST Example.

- <https://www.tensorflow.org/versions/r0.11/tutorials/mnist/beginners/index.html#mnist-for-ml-beginners>

Convolutional Neural Networks.

- <https://www.tensorflow.org/versions/r0.11/tutorials/mnist/pros/index.html#deep-mnist-for-experts>

