


RECOMMENDATION



EXAMPLES

- Diapers and beer (legend?)
- Target pregnant teenager



EXAMPLES

- Netflix Prize 2009



- Amazon Recommendation Engine

Customers Who Bought This Item Also Bought



[30 Rock: Seasons 1-3](#)
DVD ~ Tracy Morgan
★★★★★ (7)
\$60.49



[Desperate Housewives: The Complete Seasons 1-5](#)
DVD ~ Teri Hatcher
★★★★☆ (2)
\$179.99

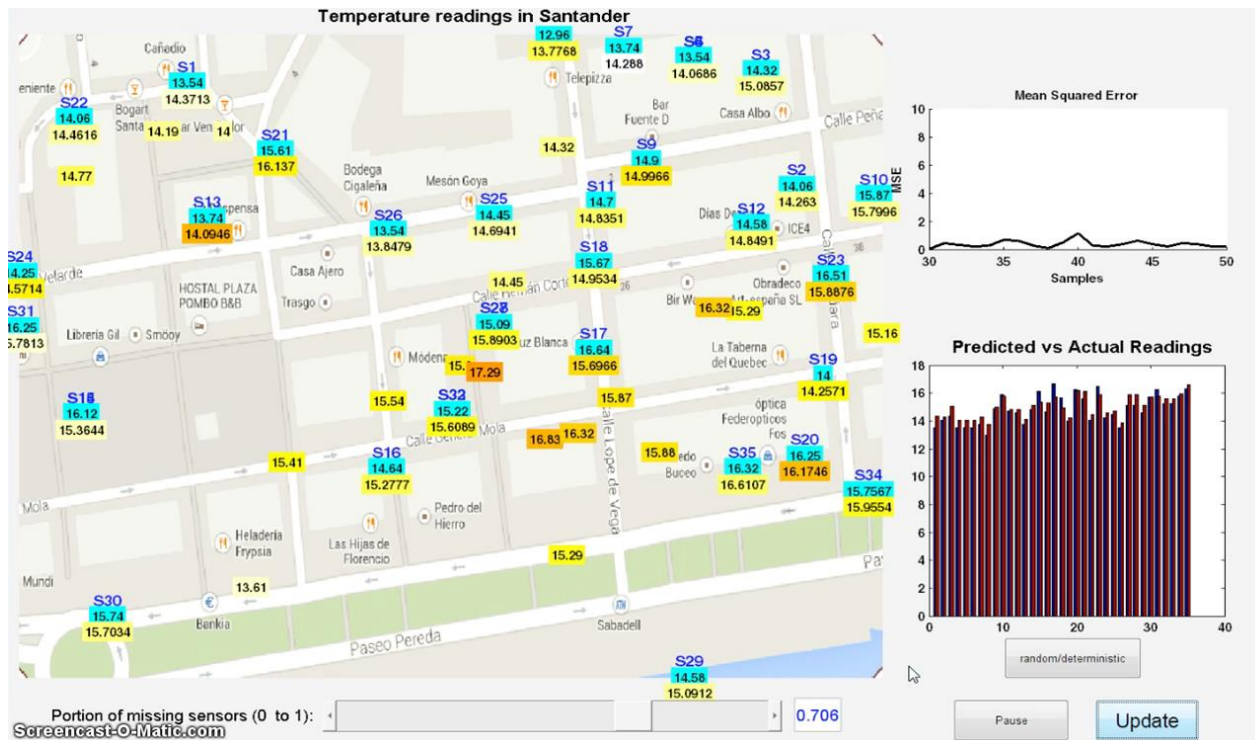


[Scrubs: The Complete Seasons 1-8](#)
DVD ~ Zach Braff
★★★★★ (2)
\$148.49



EXAMPLES

- Missing sensor data



COLLABORATIVE FILTERING

- 400,000 users
17,000 movies
But only a few ratings (1%)
- User a
Movie i
Rating $Y_{ai} \in \mathbb{R}$ (e.g. 1-5)
- Training data is the incomplete matrix Y
- Goal is to predict unobserved ratings

m movies

n users

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 5 | 5 | | | | | | 5 | | |
| | | 3 | 5 | 1 | 3 | 4 | 4 | | 4 |
| | 4 | 2 | | | 2 | | | | |
| | | 5 | | | | | | | 5 |
| 4 | 5 | | | | | | | 4 | |
| 4 | | | | | | | 4 | | |
| 5 | | 4 | 5 | 1 | | 4 | | | |
| | 4 | | | | | | | | |
| 5 | | | | 4 | | | | | |
| 5 | | | | | | 4 | | | |
| | | 5 | | | | 5 | | 3 | |

Y_{ai}



COLLABORATIVE FILTERING

- **Matrix or tensor completion problems**
- Collaborative: cross-users
Filtering: prediction
- Dimensionality reduction

A tensor is a multi-dimensional array.
e.g. $n \times m$ 2-tensor
e.g. $p \times q \times r$ 3-tensor
(2-tensor = matrix)

n users

m movies

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 5 | 5 | | | | | | 5 | | |
| | | 3 | 5 | 1 | 3 | 4 | 4 | | 4 |
| | 4 | 2 | | | 2 | | | | |
| | | 5 | | | | | | | 5 |
| 4 | 5 | | | | | | | 4 | |
| 4 | | | | | | | 4 | | |
| 5 | | 4 | 5 | 1 | | 4 | | | |
| | 4 | | | | | | | | |
| 5 | | | | 4 | | | | | |
| 5 | | | | | | 4 | | | |
| | | 5 | | | | 5 | | 3 | |

Y_{ai}



SUPERVISED OR UNSUPERVISED

Supervised Learning (Treat Y_{ai} as responses)

- Given training data of the form $((a, i), Y_{ai})$,
find a function $f: \{1, \dots, n\} \times \{1, \dots, m\} \rightarrow \mathbb{R}$.

Unsupervised Learning (Treat Y_{ai} as features)

- Given (incomplete) user ratings $Y_a \in \mathbb{R}^m$,
find structure in the data to predict missing values.

Recommendation is a *task* for which you can use either supervised or unsupervised learning algorithms.






K-NEAREST NEIGHBORS



K-NEAREST NEIGHBORS

Main Idea.


- Find **a few users** b_1, \dots, b_k (neighbors) that are similar to user a
- Use information from users b_1, \dots, b_k to predict ratings of user a



CORRELATION COEFFICIENT

Cosine Similarity

$$\cos(x, y) = \frac{x^T y}{\|x\| \|y\|}$$



Correlation Coefficient

$$\text{corr}(x, y) = \cos(x - \bar{x}, y - \bar{y}) = \frac{(x - \bar{x})^T (y - \bar{y})}{\|x - \bar{x}\| \|y - \bar{y}\|}$$

where \bar{x}, \bar{y} are the averages of the entries of x, y .



USER SIMILARITY

To compute the similarity between users a and b ,

1. Find $CR(a, b)$, the set of movies rated by both a and b .
2. Let Z_a, Z_b be the vector of ratings in $CR(a, b)$ for each user.
3. Compute the correlation coefficient between Z_a and Z_b .

$$\text{sim}(a, b) = \text{corr}(Z_a, Z_b) \in [-1, 1]$$

CR = “Common Ratings”



WEIGHTED PREDICTION

Let \bar{Y}_b denote the average of movie ratings by a user b .

To predict the rating Y_{ai} of a user a for a movie i ,

1. Rank users b who rated movie i according to value of $\text{sim}(a, b)$.
2. Let $\text{kNN}(a, i)$ be the set of highest k users.
3. Let $(Y_{ai} - \bar{Y}_a)$ be weighted average of $(Y_{bi} - \bar{Y}_b)$, $b \in \text{kNN}(a, i)$.
4. Let weight for $(Y_{bi} - \bar{Y}_b)$ be proportional to $\text{sim}(a, b)$.



ANTI-CORRELATED USERS

If $(Y_{bi} - \bar{Y}_b)$ is strongly anti-correlated with $(Y_{ai} - \bar{Y}_a)$, i.e.

$$\text{sim}(a, b) \ll 0,$$

then $-(Y_{bi} - \bar{Y}_b)$ is strongly correlated with $(Y_{ai} - \bar{Y}_a)$.

We should exploit this information because ratings are rare.



WEIGHTED PREDICTION

Let \bar{Y}_b denote the average of movie ratings by a user b .

To predict the rating Y_{ai} of a user a for a movie i ,

1. Rank users b who rated movie i according to value of $|\text{sim}(a, b)|$.
2. Let $\text{kNN}(a, i)$ be the set of highest k users.
3. Let $(Y_{ai} - \bar{Y}_a)$ be weighted average of $\pm(Y_{bi} - \bar{Y}_b)$, $b \in \text{kNN}(a, i)$.
4. Let weight for $\pm(Y_{bi} - \bar{Y}_b)$ be proportional to $|\text{sim}(a, b)|$.

$$\hat{Y}_{ai} - \bar{Y}_a = \frac{\sum_{b \in \text{kNN}(a, i)} \text{sim}(a, b)(Y_{bi} - \bar{Y}_b)}{\sum_{b \in \text{kNN}(a, i)} |\text{sim}(a, b)|}$$



DISCUSSION

$$\hat{Y}_{ai} - \bar{Y}_a = \frac{\sum_{b \in \text{kNN}(a,i)} \text{sim}(a, b)(Y_{bi} - \bar{Y}_b)}{\sum_{b \in \text{kNN}(a,i)} |\text{sim}(a, b)|}$$

- Formula is not sensitive to the bias (mean) of each user, but it is sensitive to the spread (variance) of each user.
- There is no training loss, no training algorithm for kNN.





MATRIX FACTORIZATION



SUBSPACE LEARNING


$k \ll m$ means k is much smaller than m

Main Idea.

- Completed rating vectors $\hat{Y}_1, \dots, \hat{Y}_n \in \mathbb{R}^m$ lie in some k -dimensional subspace, $k \ll m$.
- i.e. there exists vectors $V_{:1}, \dots, V_{:k} \in \mathbb{R}^m$ such that for all users a ,

$$\hat{Y}_a = U_{a1} V_{:1} + U_{a2} V_{:2} + \dots + U_{ak} V_{:k}$$

for some coefficients $U_{a1}, \dots, U_{ak} \in \mathbb{R}$.



MATRIX FACTORIZATION

$$\begin{aligned}\hat{Y}_1^\top &= U_{11} V_{:1}^\top + \cdots + U_{1k} V_{:k}^\top \\ \hat{Y}_2^\top &= U_{21} V_{:1}^\top + \cdots + U_{2k} V_{:k}^\top \\ &\vdots \\ \hat{Y}_n^\top &= U_{n1} V_{:1}^\top + \cdots + U_{nk} V_{:k}^\top\end{aligned} \quad \Rightarrow \quad \begin{aligned}\hat{Y} &= U_{:1} V_{:1}^\top + \cdots + U_{:k} V_{:k}^\top \\ &= (U_{:1}, \dots, U_{:k}) \begin{pmatrix} V_{:1}^\top \\ \vdots \\ V_{:k}^\top \end{pmatrix} \\ &= (U_{:1}, \dots, U_{:k})(V_{:1}, \dots, V_{:k})^\top \\ &= UV^\top\end{aligned}$$

where $U_{:1}, \dots, U_{:k}$ are the columns of $U \in \mathbb{R}^{n \times k}$,
and $V_{:1}, \dots, V_{:k}$ are the columns of $V \in \mathbb{R}^{m \times k}$.



MATRIX RANK

outer
product!

Definition. A matrix $Y \in \mathbb{R}^{n \times m}$ is of rank-one if $Y = ab^\top$ for some non-zero vectors $a \in \mathbb{R}^n, b \in \mathbb{R}^m$.

Definition. The rank of a matrix Y is the smallest number k of rank-one matrices required in the sum

$$Y = Y^{(1)} + \dots + Y^{(k)}.$$

If $Y = UV^\top$ with $U \in \mathbb{R}^{n \times k}, V \in \mathbb{R}^{m \times k}$, i.e.

$$Y = UV^\top = U_{:1} V_{:1}^\top + U_{:2} V_{:2}^\top + \dots + U_{:k} V_{:k}^\top,$$


then the rank Y is at most k .



LOW-RANK APPROXIMATION

Main Idea.

- Find low-rank *complete* matrix \hat{Y} that is closest to *incomplete* matrix Y .
- \hat{Y} is called a *low-rank approximation* of Y .



PREDICTION

For an unknown rating Y_{ai} of user a for movie i , we predict

$$\hat{Y}_{ai} = (UV^T)_{ai} = U_a^T V_i$$

where U_a is the a -th row of U and V_i is the i -th row of V .

Compare this with

$$\hat{Y} = U_{:1} V_{:1}^T + U_{:2} V_{:2}^T + \dots + U_{:k} V_{:k}^T$$

where $U_{:j}$ is the j -th column of U and $V_{:j}$ is the j -th column of V .



TRAINING LOSS

We regularize
because there
are infinitely
many solutions
 $U' = cU, V' = c^{-1}V$

Apply squared loss to each observed rating Y_{ai} .

$$\begin{aligned}\mathcal{L}_{n,k,\lambda}(U, V; Y) &= \sum_{(a,i) \in D} \frac{1}{2} (Y_{ai} - (UV^\top)_{ai})^2 + \frac{\lambda}{2} \|U\|^2 + \frac{\lambda}{2} \|V\|^2 \\ &= \sum_{(a,i) \in D} \frac{1}{2} (Y_{ai} - U_a^\top V_i)^2 + \frac{\lambda}{2} \sum_a \|U_a\|^2 + \frac{\lambda}{2} \sum_i \|V_i\|^2\end{aligned}$$

where D is the set of (a, i) where Y_{ai} is observed.

Definition. The *Frobenius* norm $\|U\|$ of a matrix $U \in \mathbb{R}^{n \times k}$ is

$$\|U\| = \sqrt{\sum_{a=1}^n \sum_{b=1}^k U_{ab}^2} .$$



ALTERNATING LEAST SQUARES

Coordinate Descent (optimization).

Repeat until convergence:

1. Fix V and minimize $\mathcal{L}_{n,k,\lambda}(U, V; Y)$ over U .
2. Fix U and minimize $\mathcal{L}_{n,k,\lambda}(U, V; Y)$ over V .



ALTERNATING LEAST SQUARES

1. Initialize $V_1, V_2, \dots, V_m \in \mathbb{R}^k$ randomly.
2. Repeat until convergence:
 - a. For each user a , find U_a that minimizes
$$\sum_{i: (a,i) \in D} \frac{1}{2} (Y_{ai} - U_a^\top V_i)^2 + \frac{\lambda}{2} \|U_a\|^2$$
 - b. For each movie i , find V_i that minimizes
$$\sum_{a: (a,i) \in D} \frac{1}{2} (Y_{ai} - U_a^\top V_i)^2 + \frac{\lambda}{2} \|V_i\|^2$$

These are
standard
linear
regression
problems.



DISCUSSION

User Bias

- Do we subtract the average ratings of each user?

Optimization.

- Like k-means, the algorithm converges to a local minimum.
- Perform multiple initializations, and pick best result.

Generalization.

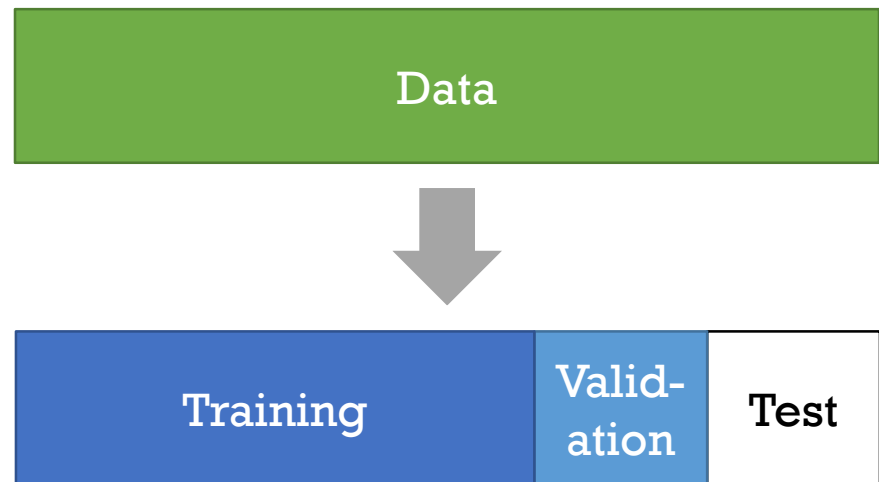
- Use **validation** to pick right **hyperparameters** k and λ .



VALIDATION SET

Split the data into

- **Test set** \mathcal{S}_*
For evaluating, reporting performance at the end
- **Training set** \mathcal{S}_n
For training optimal parameters in a model
- **Validation set** \mathcal{S}_{val}
For model selection, e.g. picking k in k-means, picking λ in ridge regression. Acts as a proxy for test set.



VALIDATION LOSS

The *validation error* is the test loss applied to the validation set.

The *training error* is the test loss applied to the training set, and it may be different from the training loss used for optimization.

Example. Ridge Regression

Test loss/error $\mathcal{R}(\hat{\theta}; \mathcal{S}_*) = \frac{1}{n} \sum_{(x,y) \in \mathcal{S}_*} \frac{1}{2} (y - \hat{\theta}^\top x)^2$


Validation loss/error $\mathcal{R}(\hat{\theta}; \mathcal{S}_{\text{val}}) = \frac{1}{n} \sum_{(x,y) \in \mathcal{S}_{\text{val}}} \frac{1}{2} (y - \hat{\theta}^\top x)^2$

Training error $\mathcal{R}(\hat{\theta}; \mathcal{S}_n) = \frac{1}{n} \sum_{(x,y) \in \mathcal{S}_n} \frac{1}{2} (y - \hat{\theta}^\top x)^2$

Training loss $\mathcal{L}_{n,\lambda}(\theta; \mathcal{S}_n) = \frac{1}{n} \sum_{\text{data } (x,y)} \frac{1}{2} (y - \theta^\top x)^2 + \frac{\lambda}{2} \|\theta\|^2$

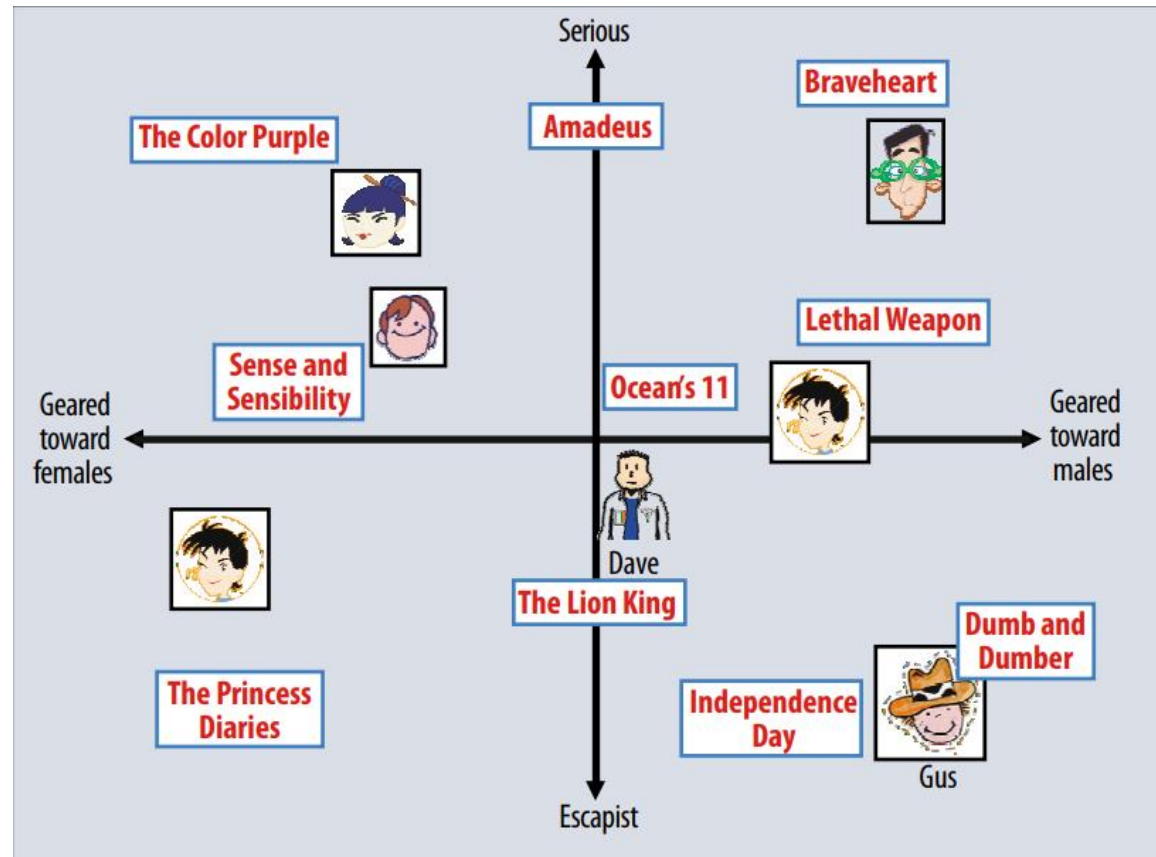


MODEL SELECTION




NETFLIX RESULTS

Plot each movie i according to first two coordinates of $V_i \in \mathbb{R}^k$, i.e. V_{i1} , V_{i2} .



NETFLIX RESULTS

Winning team
*Bellkor's
Pragmatic Chaos*
essentially used
a matrix
factorization
model with
additional
parameters
for biases,
user attributes
and time
dependencies.



SUMMARY

- Matrix Completion
 - Recommender Systems
 - Collaborative Filtering
- k -Nearest Neighbors
 - Correlation Coefficient
 - Weighted Prediction
- Matrix Factorization
 - Subspace Learning
 - Low-rank Approximation
 - Regularized Training Loss
 - Alternating Least Squares



INTENDED LEARNING OUTCOMES

Matrix Completion

- Recognize that the key problem in collaborative filtering is matrix completion. Give examples of collaborative filtering.
- Explain how dimensionality reduction helps matrix completion.

k-Nearest Neighbors

- Compute the user similarity using the correlation coefficient.
- Compute the k-nearest neighbors ranked by user similarity.
- Predict an unknown rating using weighted prediction.



INTENDED LEARNING OUTCOMES

Matrix Factorization

- Describe the relationship between subspace learning, low-rank approximation and matrix factorization.
- Write down the regularized training loss. Explain why regularization is necessary for matrix factorization.
- Describe the alternating least squares algorithm, and explain why it minimizes the training loss. Apply the algorithm in a recommendation problem.
- Describe how validation can be used to select suitable hyperparameters k and λ for good generalization.

