

Generative Model

Autoencoder (AE) and Generative adversarial network (GAN)

Weijian Deng

Overview

- Thanks for giving me the opportunity to talk about the generative model
- This talk covers
 - Supervised learning vs. unsupervised learning
 - Generative model
 - Autoencoder (sparse AE); generative adversarial network (GAN)

Supervised learning vs. Unsupervised learning

Supervised learning

Data: (x, y)

X is sample, y is label

Goal: learn a function to map $x \rightarrow y$

Image classification



This image is CC0 public domain

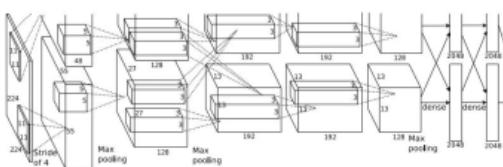


Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

Vector:
4096

Fully-Connected:
4096 to 1000

Class Scores
Cat: 0.9
Dog: 0.05
Car: 0.01
...

cat

Supervised learning vs. Unsupervised learning

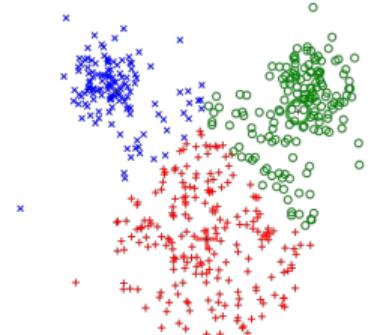
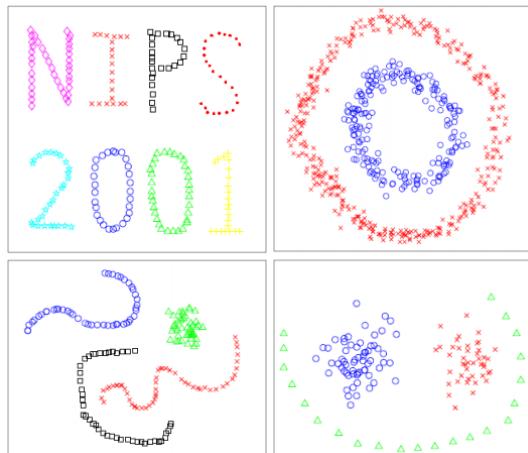
Unsupervised learning

Data: x

X is sample, no labels

Goal: learn underlying hidden structure of data

Clustering



Supervised learning vs. Unsupervised learning

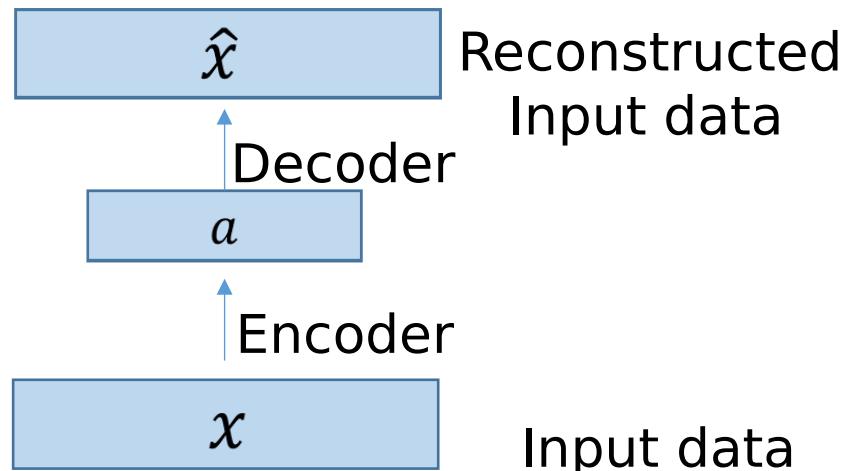
Unsupervised learning

Data: x

X is sample, no labels

Goal: learn underlying hidden structure of data

Autoencoder



Loss function: $\|x - \hat{x}\|^2$



Supervised learning vs. Unsupervised learning

Supervised learning

Data: (x, y)

X is sample, y is label

Example: image classification

Goal: learn a function to map $x \rightarrow y$

Unsupervised learning

Data: x

X is sample, **no label**

Example: clustering, Autoencoder

Goal: learn underlying hidden

structure of data

Supervised learning vs. Unsupervised learning

Supervised learning

Data: (x, y)

X is sample, y is label

Example: image classification

Goal: learn a function to map $x \rightarrow y$

Unsupervised learning *understand structure of*

Data: x

X is sample, **no label**

Example: clustering, Autoencoder

Goal: learn underlying hidden

visual world

structure of data

Generative model

*What I cannot create, I do not understand
—Richard Feynman*

Generative model



Real images distribution

$$P_{data}(x)$$

Real images distribution

generated images distribution

$$P_G(x)$$

generated images distribution

We want to learn $P_G(x)$ similar to $P_{data}(x)$

Generative model



Real images distribution

Real images distribution
 $P_{data}(x)$

generated images distribution

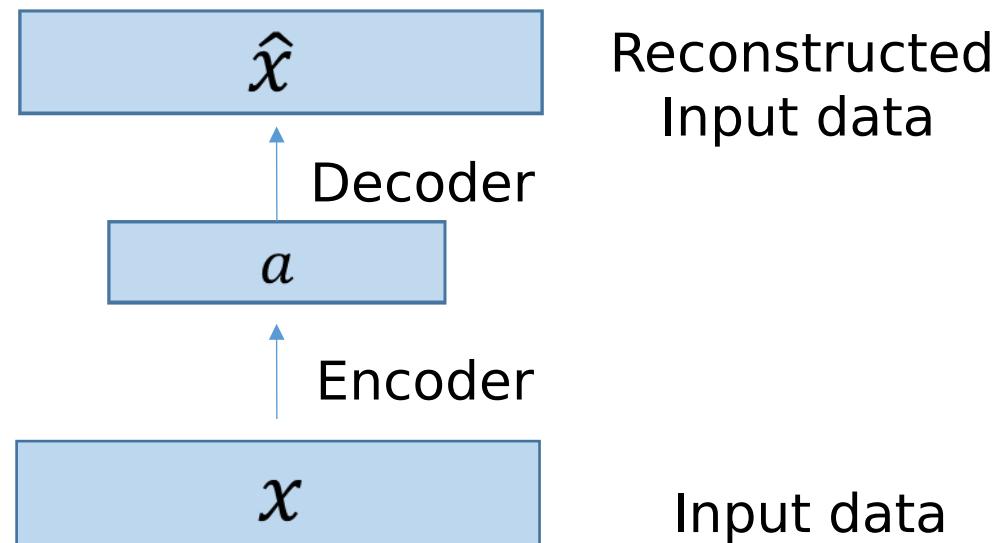
generated images distribution
 $P_G(x)$

How to learn $P_G(x)$?

Autoencoder

Autoencoder

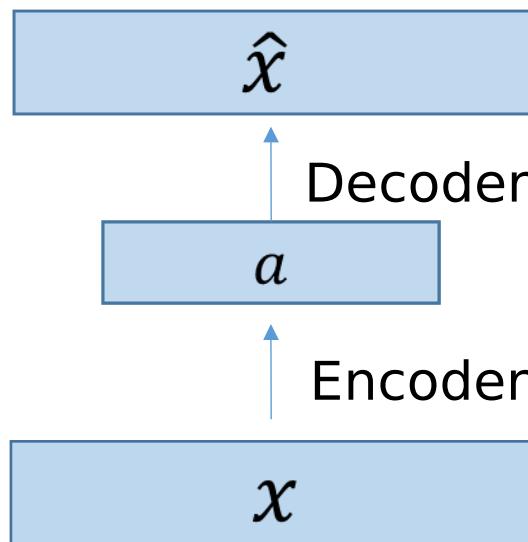
Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data



Autoencoder

Autoencoder

Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data



Reconstructed
Input data



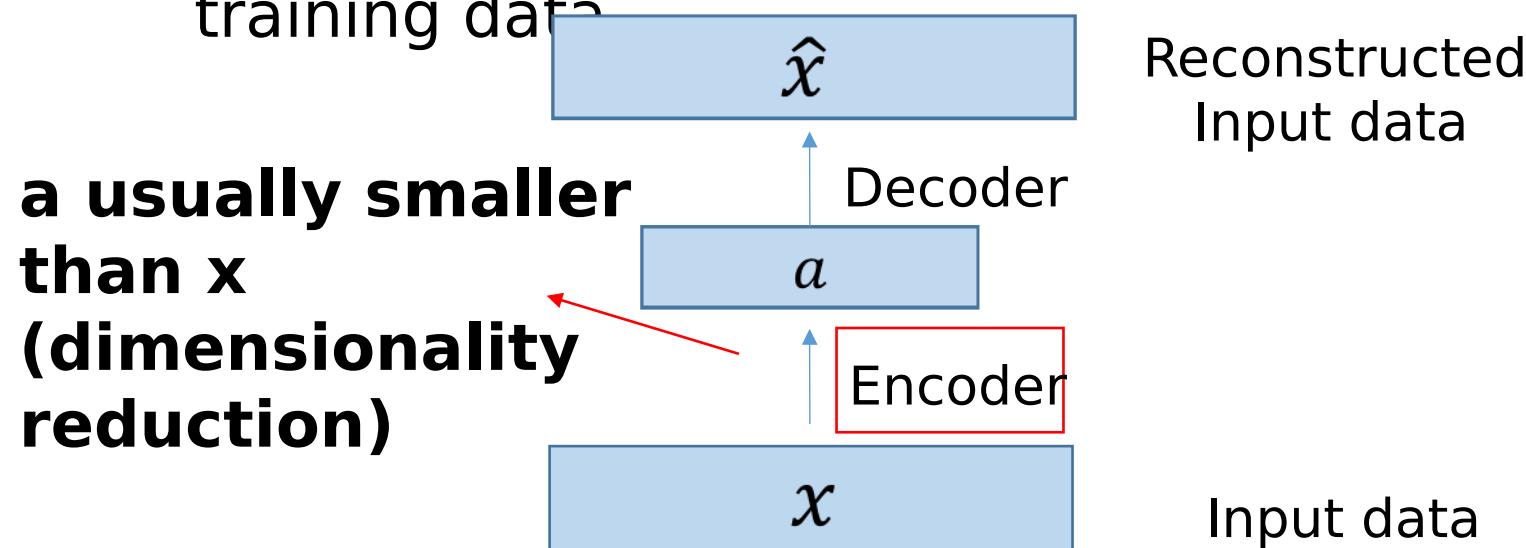
Input data



Autoencoder

Autoencoder

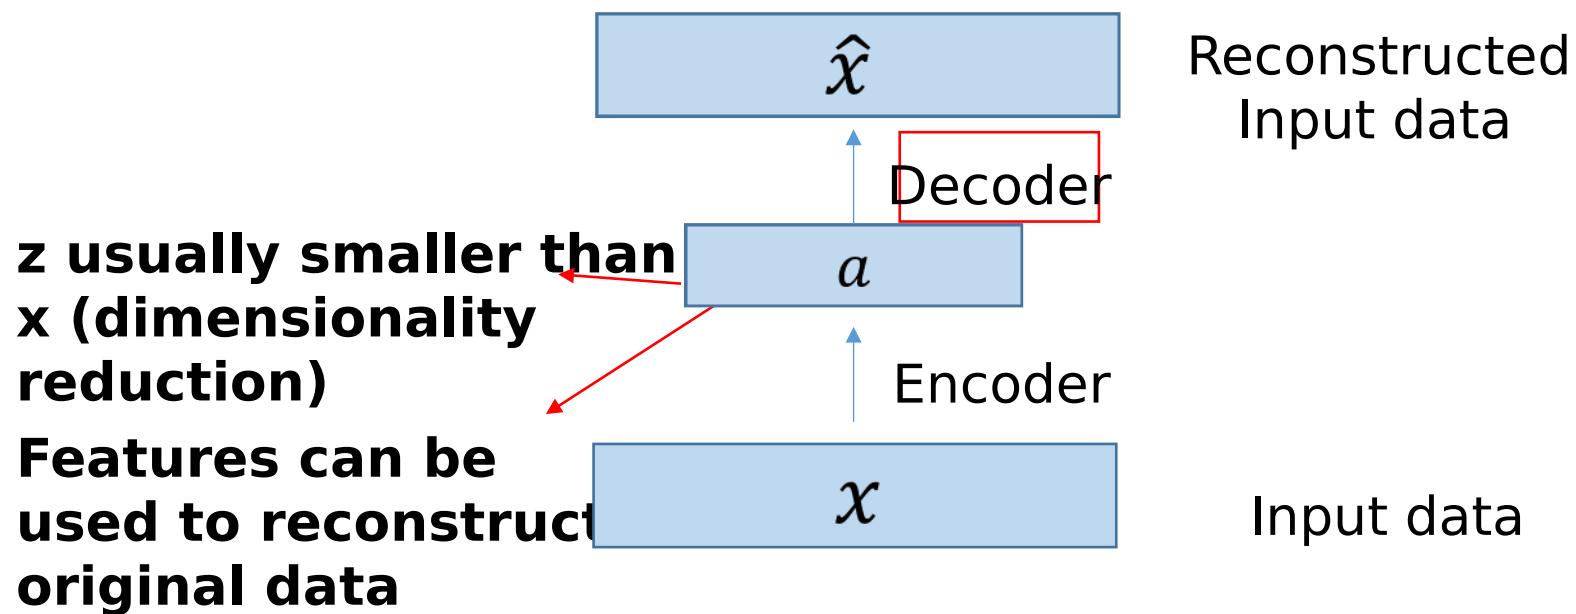
Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data



Autoencoder

Autoencoder

Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

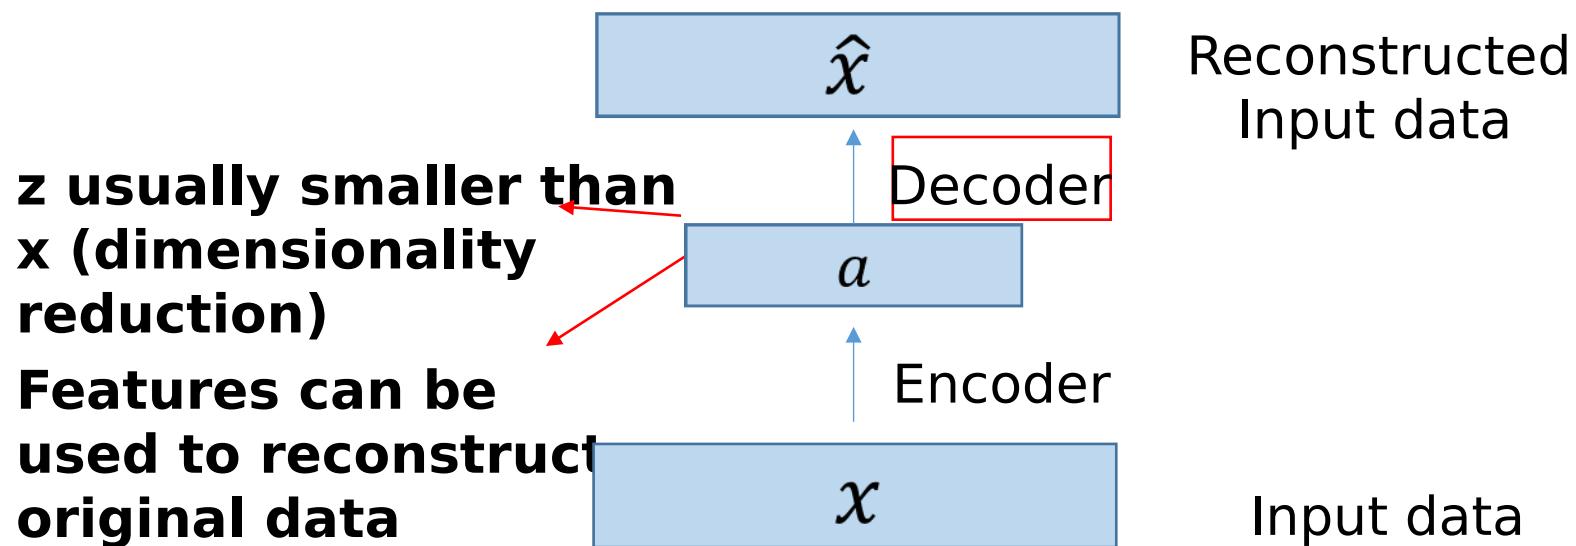


Autoencoder

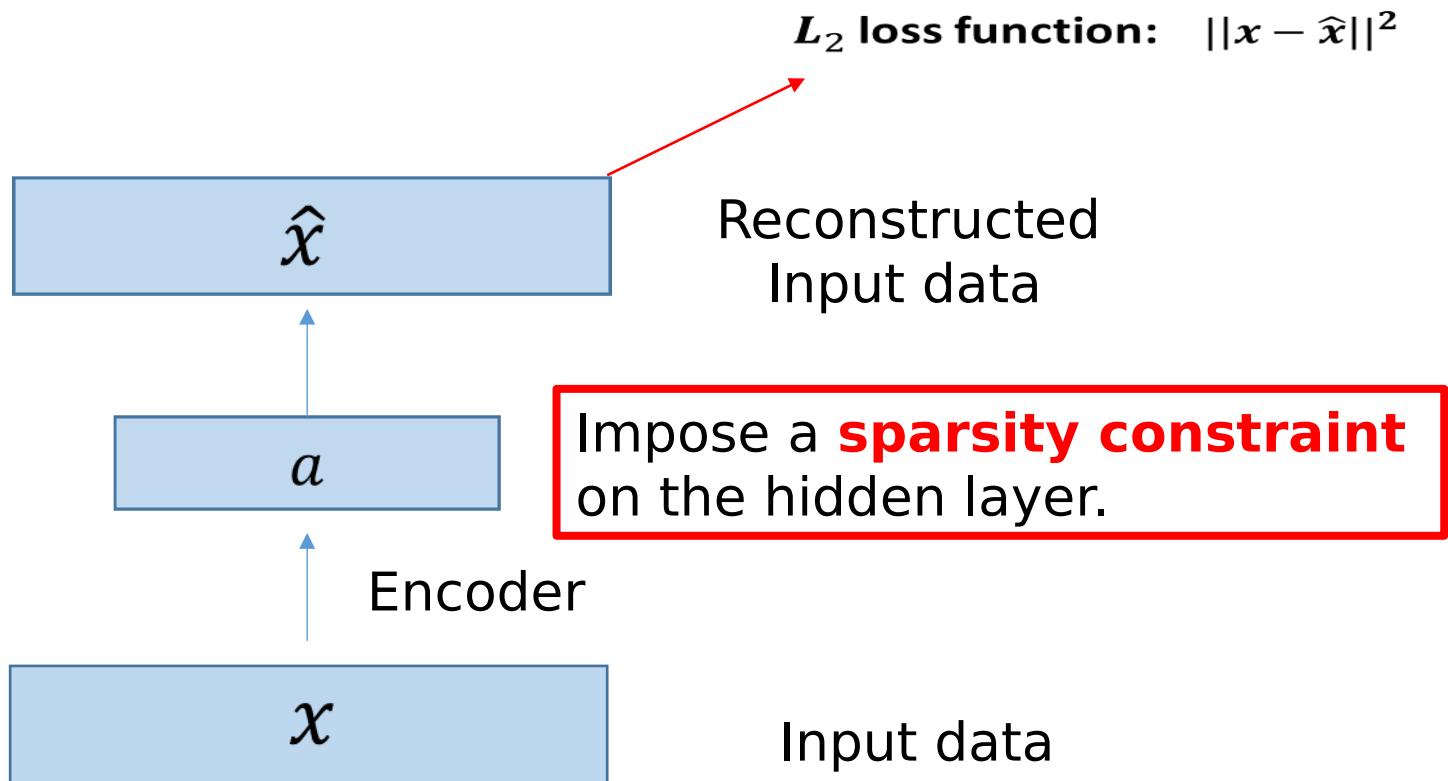
Autoencoder

Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

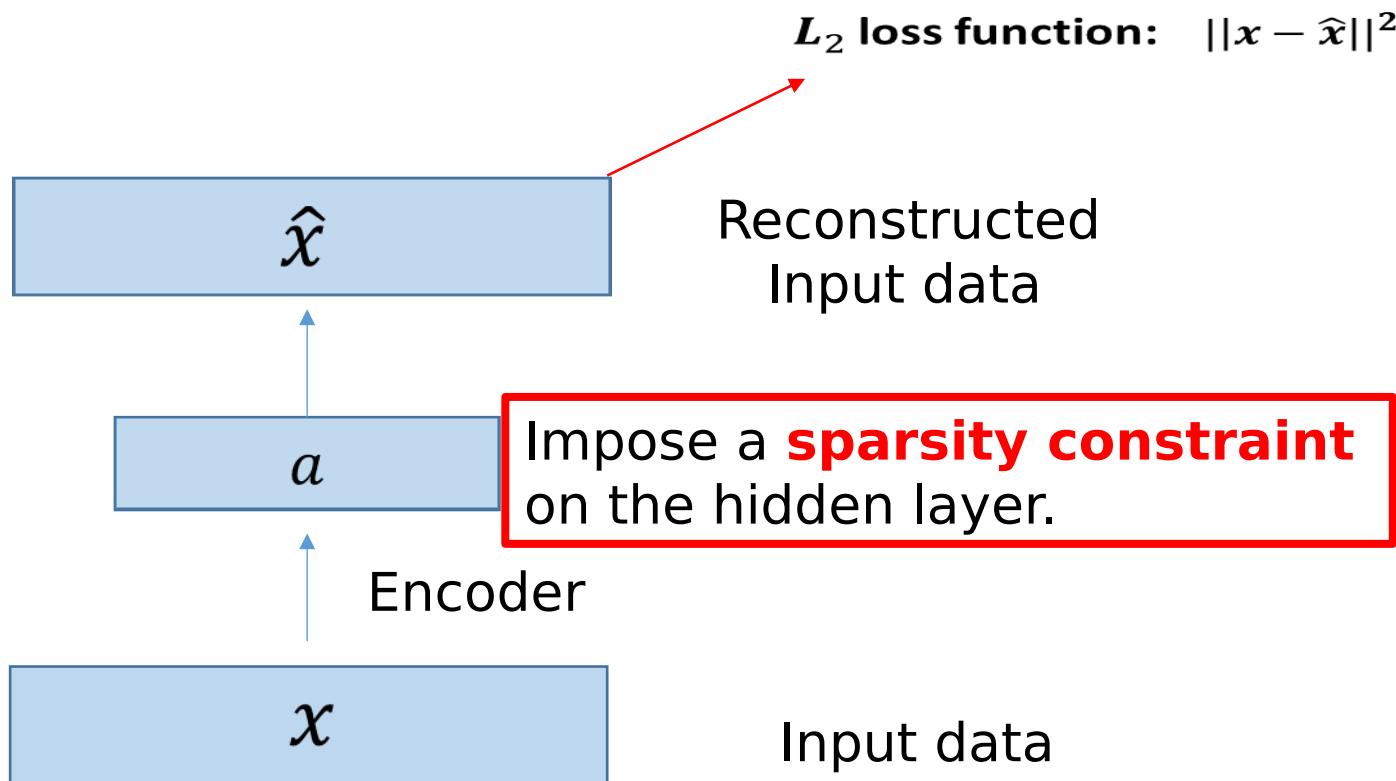
$$L_2 \text{ loss function: } ||x - \hat{x}||^2$$



Sparse Autoencoder

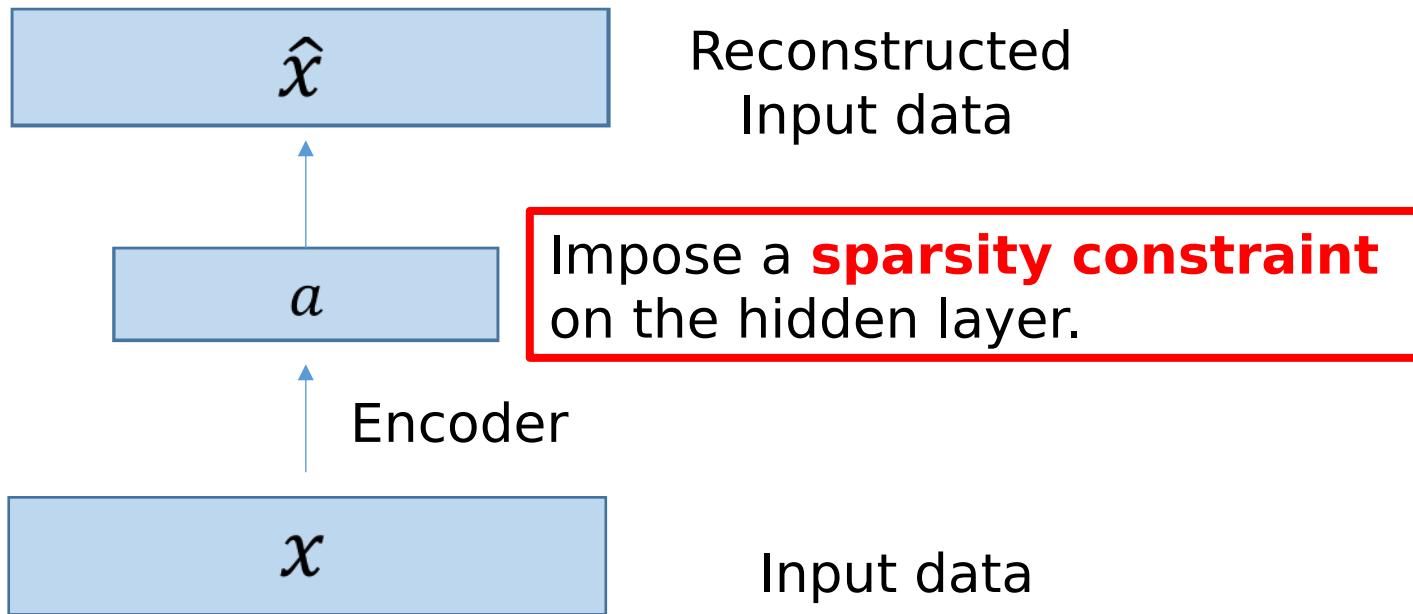


Sparse Autoencoder



We would like the **average activation** of each hidden neuron to be a small value (0.05).

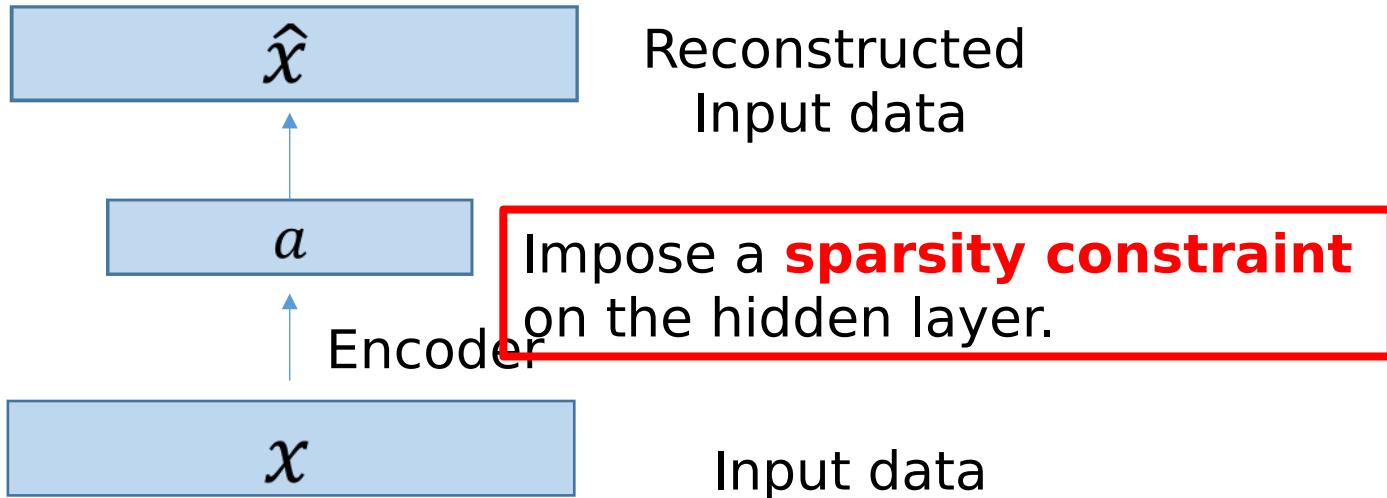
Sparse Autoencoder



$\hat{p}_j = \frac{1}{m} \sum_{i=1}^m a_j(f(x^{(i)}))$ a_j denotes the activation of hidden unit j ; we have m samples

$$a_j(f(x^{(i)})) = \text{Sigmoid}(W_j^T x^{(i)} + b)$$

Sparse Autoencoder



$\hat{p}_j = \frac{1}{m} \sum_{i=1}^m a_j(f(x^{(i)}))$ a_j denotes the activation of hidden unit j ; we have m samples

$$\sum_{j=1}^s KL(p||\hat{p}_j)$$

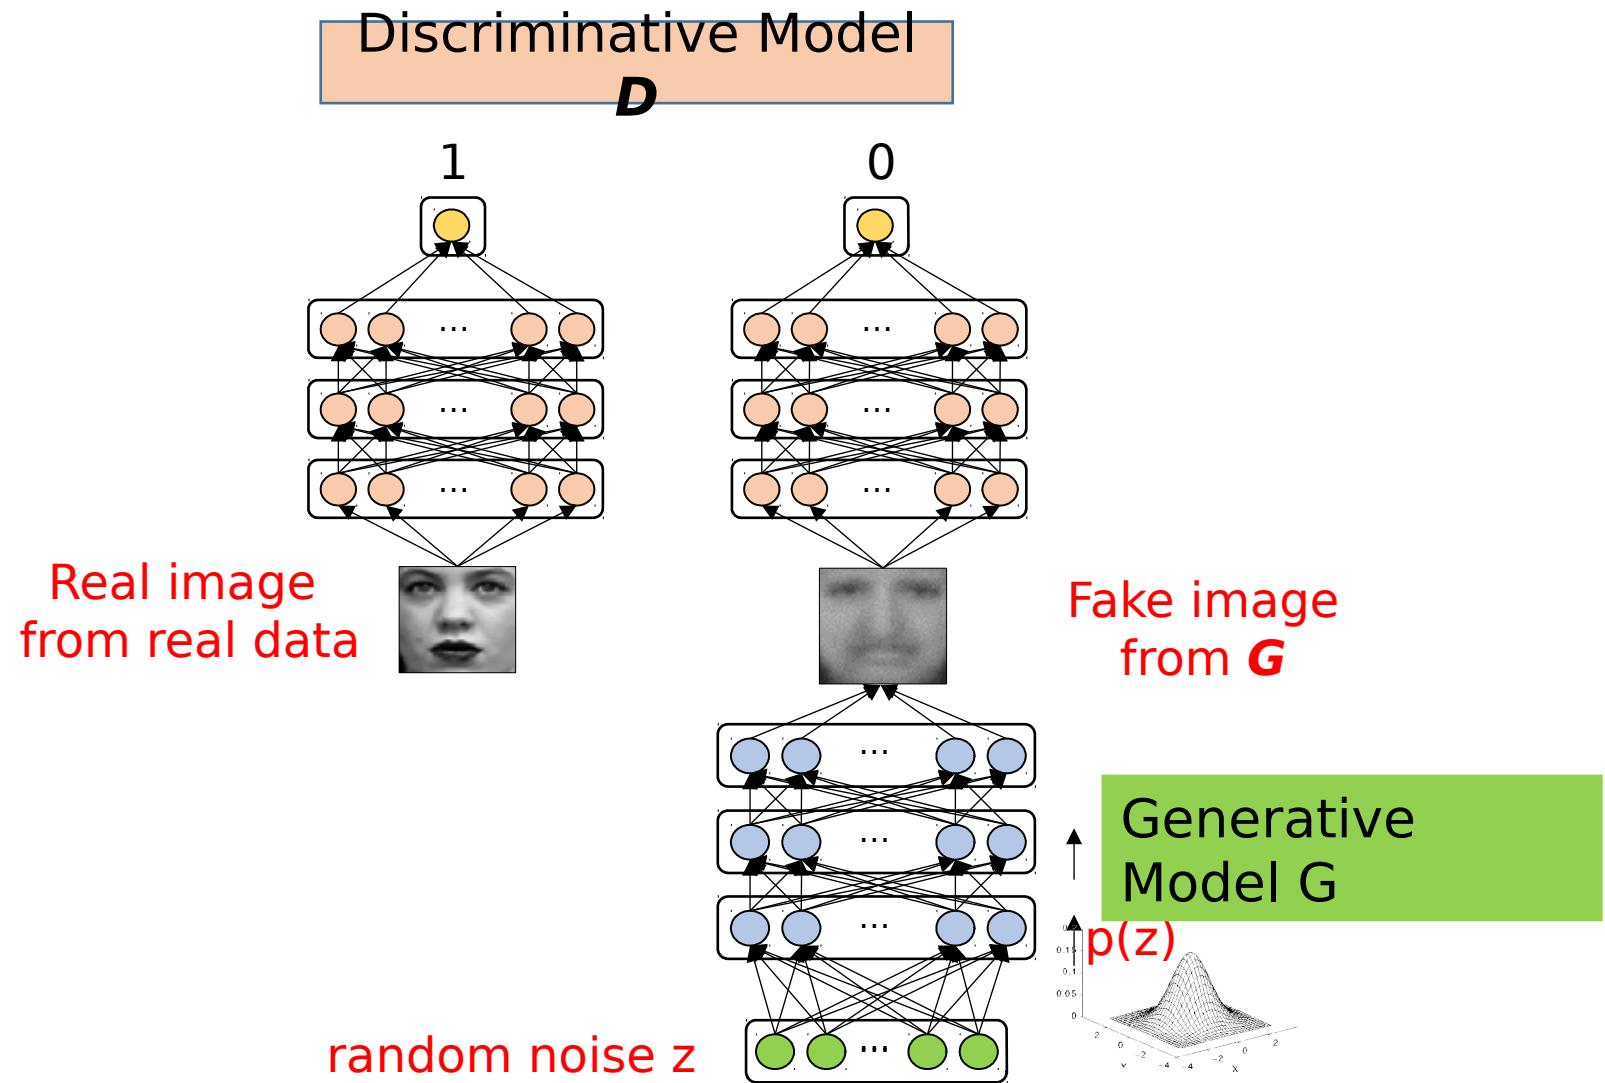
s is the number of neurons in the hidden layer; $P = 0.05$

$KL(p||\hat{p}_j) = p \log \frac{p}{\hat{p}_j} + (1-p) \log \frac{1-p}{1-\hat{p}_j}$ is the Kullback-Leibler (KL) divergence.
(KL-divergence is a standard function for measuring how different two distributions are)

Generative adversarial network (GAN)

GAN is the most important upcoming breakthrough in deep learning.
—Yan LeCun

Generative adversarial network (GAN)



Generative adversarial network (GAN)

Which of the followings are generated/ fake images ?



A



B



C



D

Generative adversarial network (GAN)

Which of the followings are generated/ fake images ?



A



B



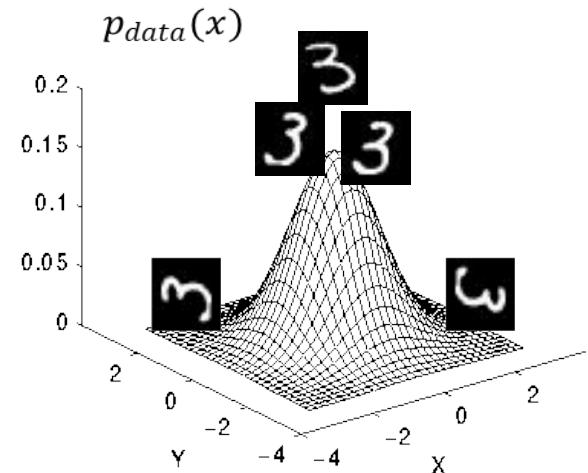
C



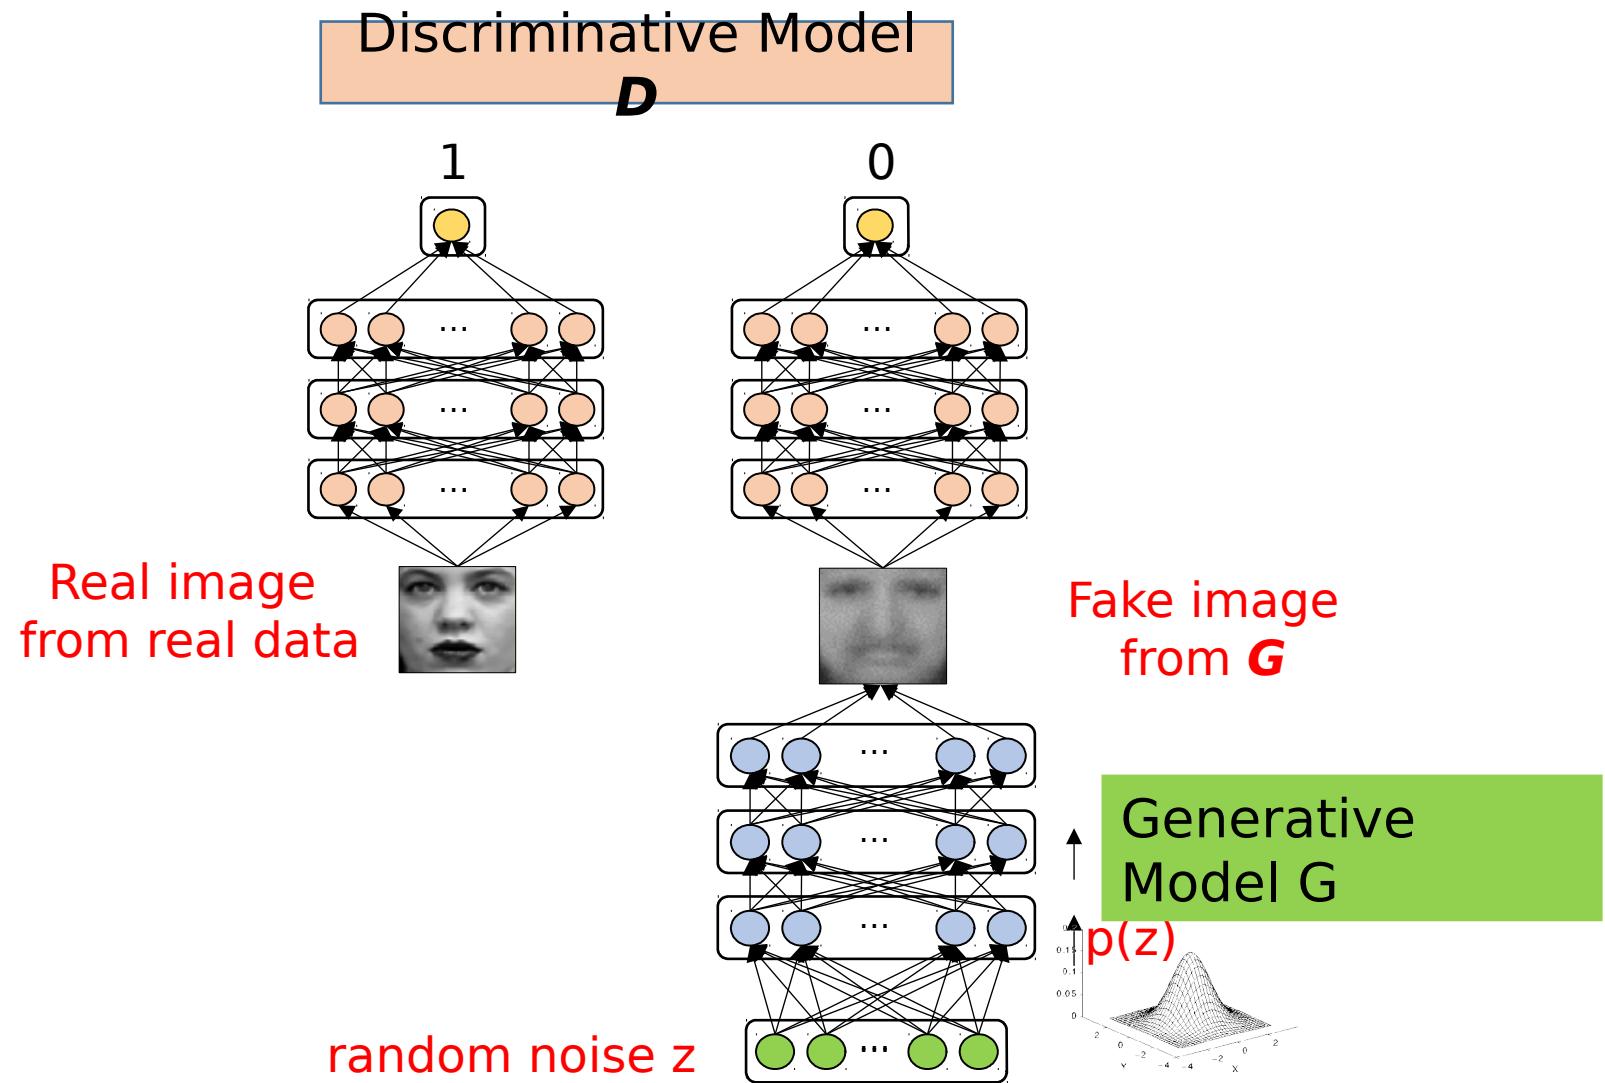
D

Generative?

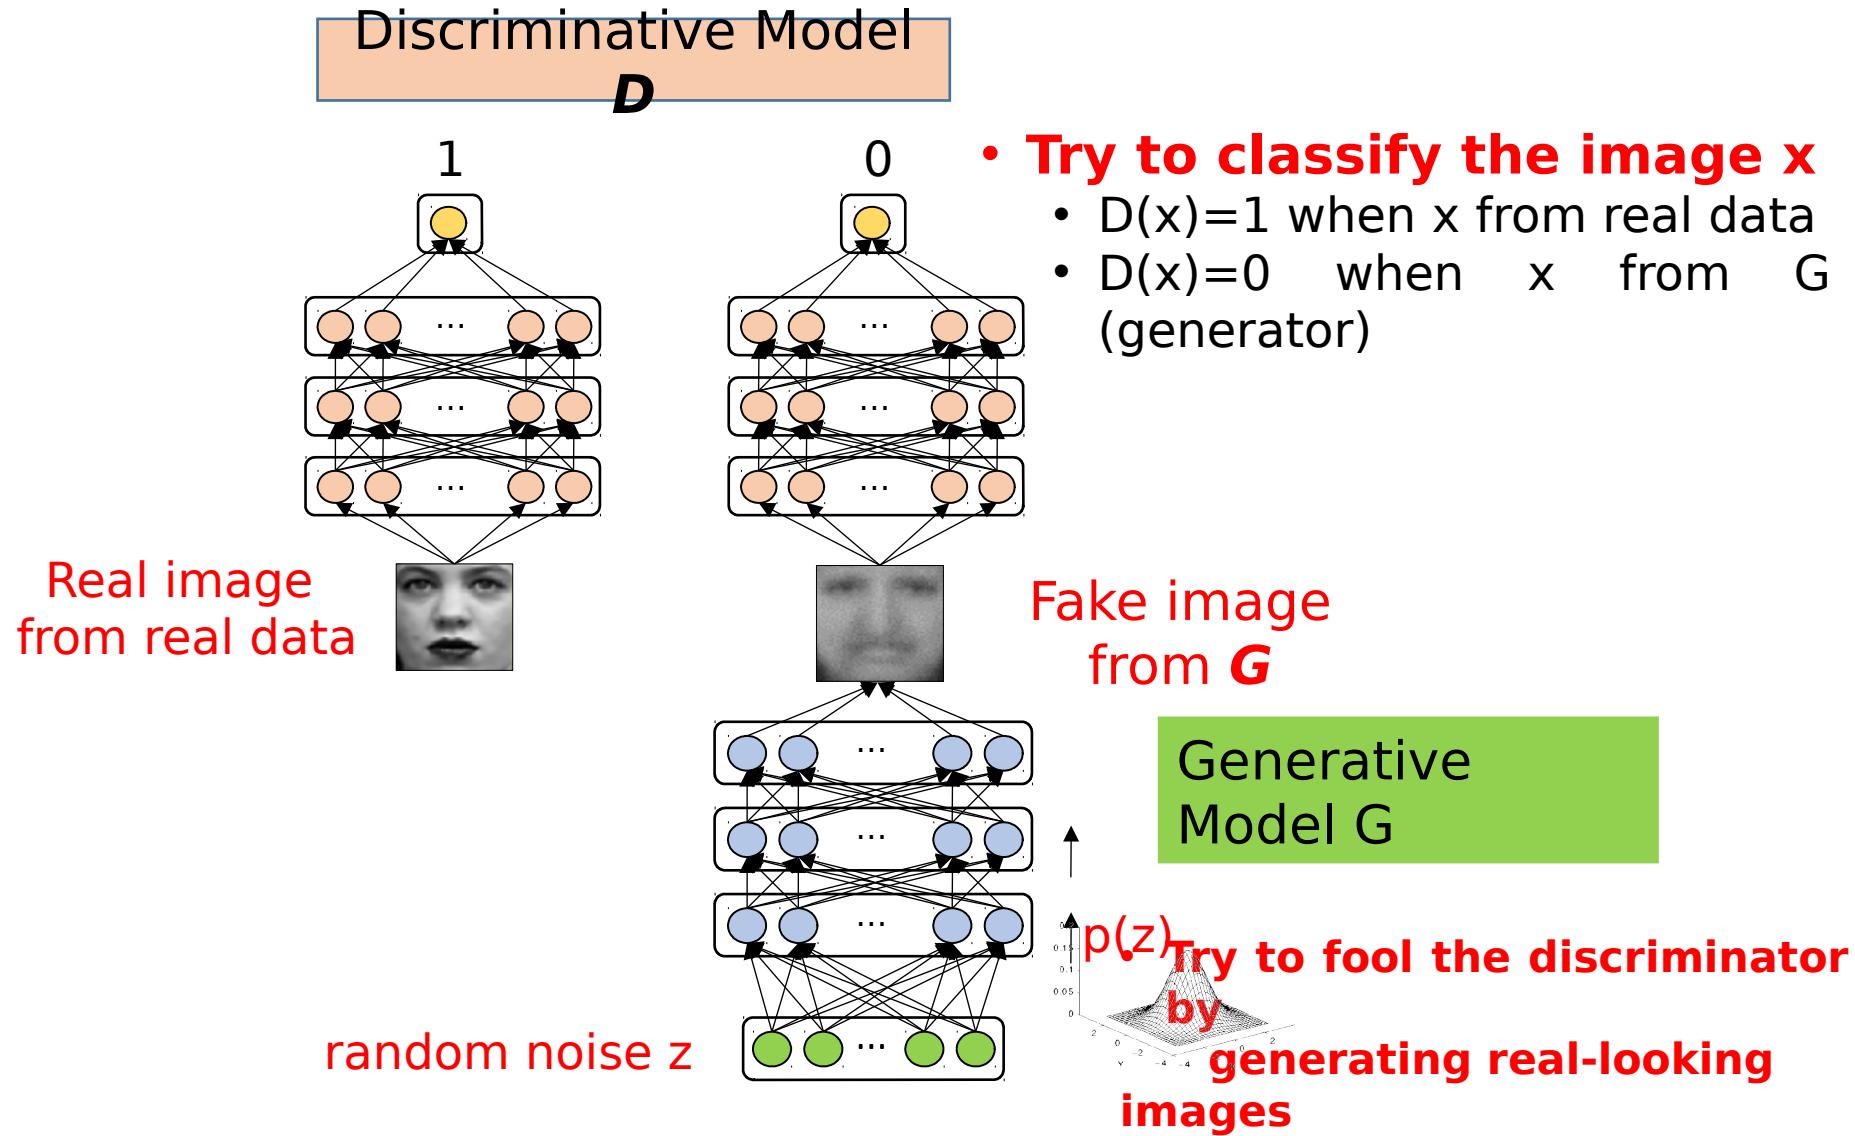
- Data: 3 3 3 3 3 3 3 3 ...
- Discriminative model: $p_D(x; \theta_D)$
- Generative model: $p_G(x; \theta_G)$
 - True data distribution: $p_{data}(x)$
 - Train $p_G(x) \approx p_{data}(x)$



Generative adversarial network (GAN)



Generative adversarial network (GAN)

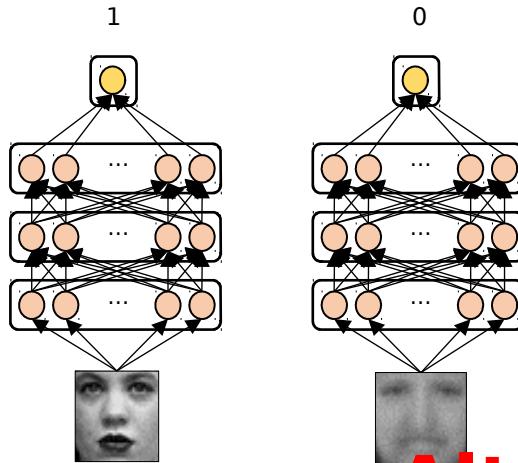


Generative adversarial network (GAN)

$$\min_G \max_D V(D, G)$$

$$V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

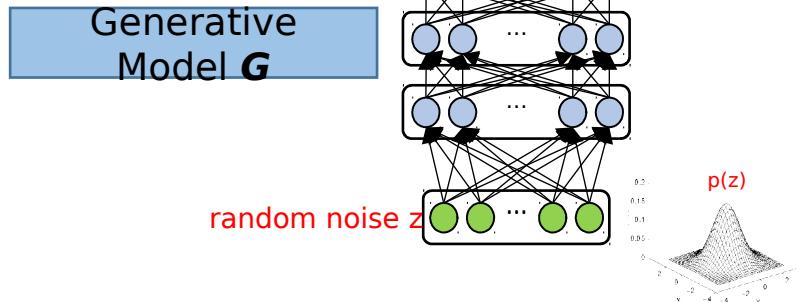
Discriminative Model **D**



Two-player game

Alternately optimize **G** and **D**

Generative Model **G**



Generative adversarial network (GAN)

$$\min_G \max_D V(D, G)$$

$$V(D, G) = \mathbb{E}_{x \sim p_{data(x)}} [\log D(x)] + \mathbb{E}_{z \sim p_{z(z)}} [\log(1 - D(G(z)))]$$

Step 1: Optimize D

- Fixed G, maximize V(D):

$$\max_D V_G(D) = \max_D \left[\mathbb{E}_{x \sim p_{data(x)}} [\log D(x)] + \mathbb{E}_{z \sim p_{z(z)}} [\log(1 - D(G(z)))] \right]$$

- From sample $x^{(i)}, z^{(i)}$

$$\max_D \left[\sum_{i=1}^m \left[\log D(x^{(i)}) + \log(1 - D(G(z^{(i)}))) \right] \right]$$

- Binary Classification (logistic loss):

- Sample from data: label=1
 - Sample from generator: label = 0

Generative adversarial network (GAN)

$$\min_G \max_D V(D, G)$$

$$V(D, G) = \mathbb{E}_{x \sim p_{data(x)}} [\log D(x)] + \mathbb{E}_{z \sim p_{z(z)}} [\log(1 - D(G(z)))]$$

Step 2: Optimize G

- Fixed D, minimize $V_D(G)$:

$$\min_G V_D(G) = \min \left[\mathbb{E}_{z \sim p_{z(z)}} [\log(1 - D(G(z)))] \right]$$

Try to make
Try to make $D(G(z)) = 1$

Generative adversarial network (GAN)

$$\min_G \max_D V(D, G)$$

$$V(D, G) = \mathbb{E}_{x \sim p_{data(x)}} [\log D(x)] + \mathbb{E}_{z \sim p_{z(z)}} [\log(1 - D(G(z)))]$$

Step 2: Optimize G

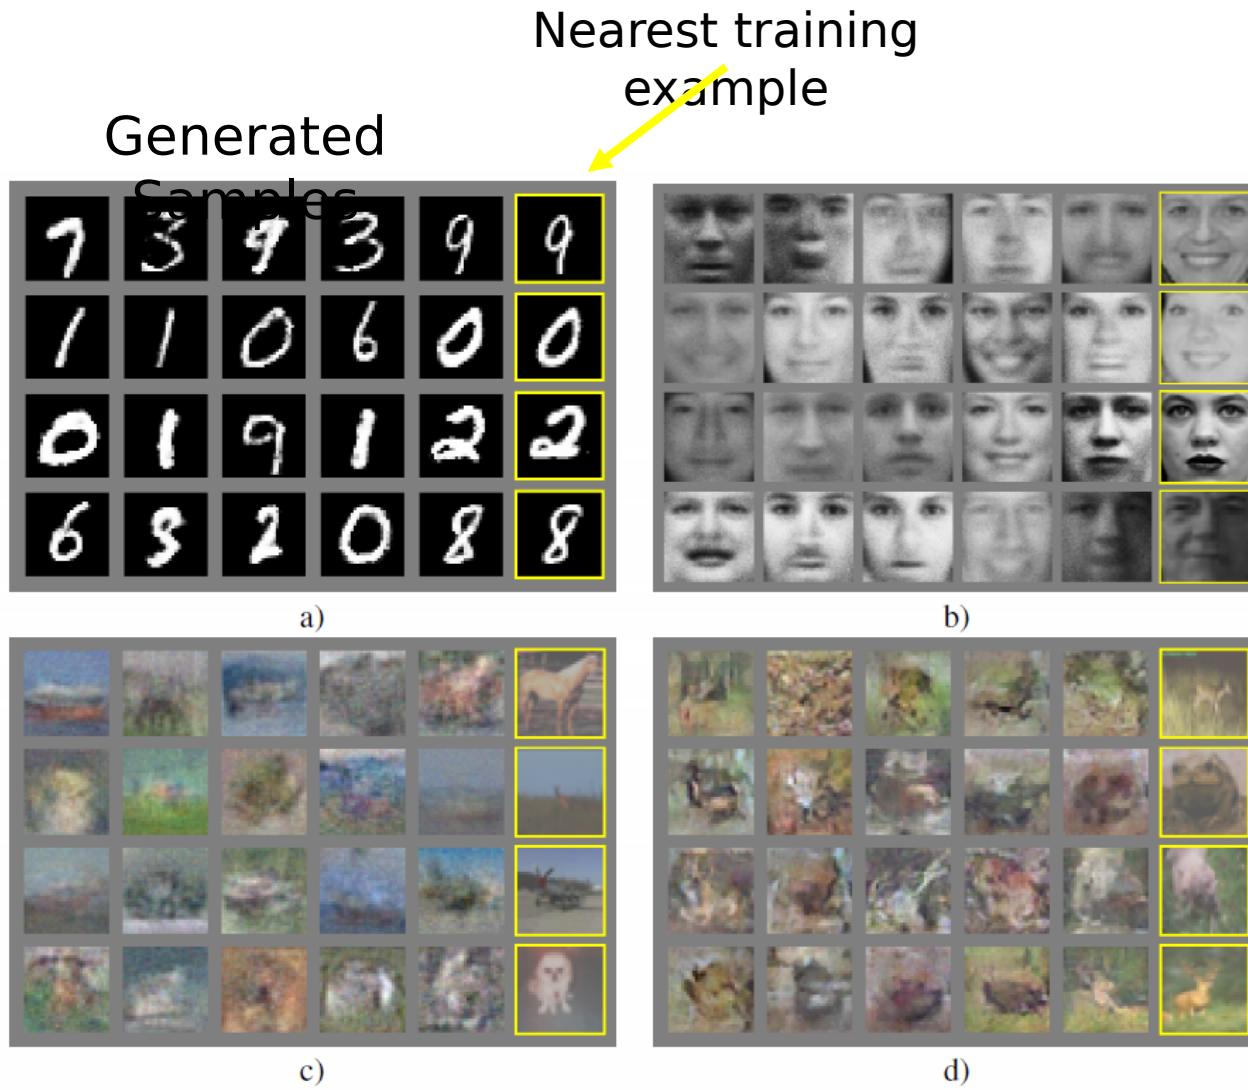
- Fixed D, minimize $V_D(G)$:

$$\min_G V_D(G) = \min \left[\mathbb{E}_{z \sim p_{z(z)}} [\log(1 - D(G(z)))] \right]$$

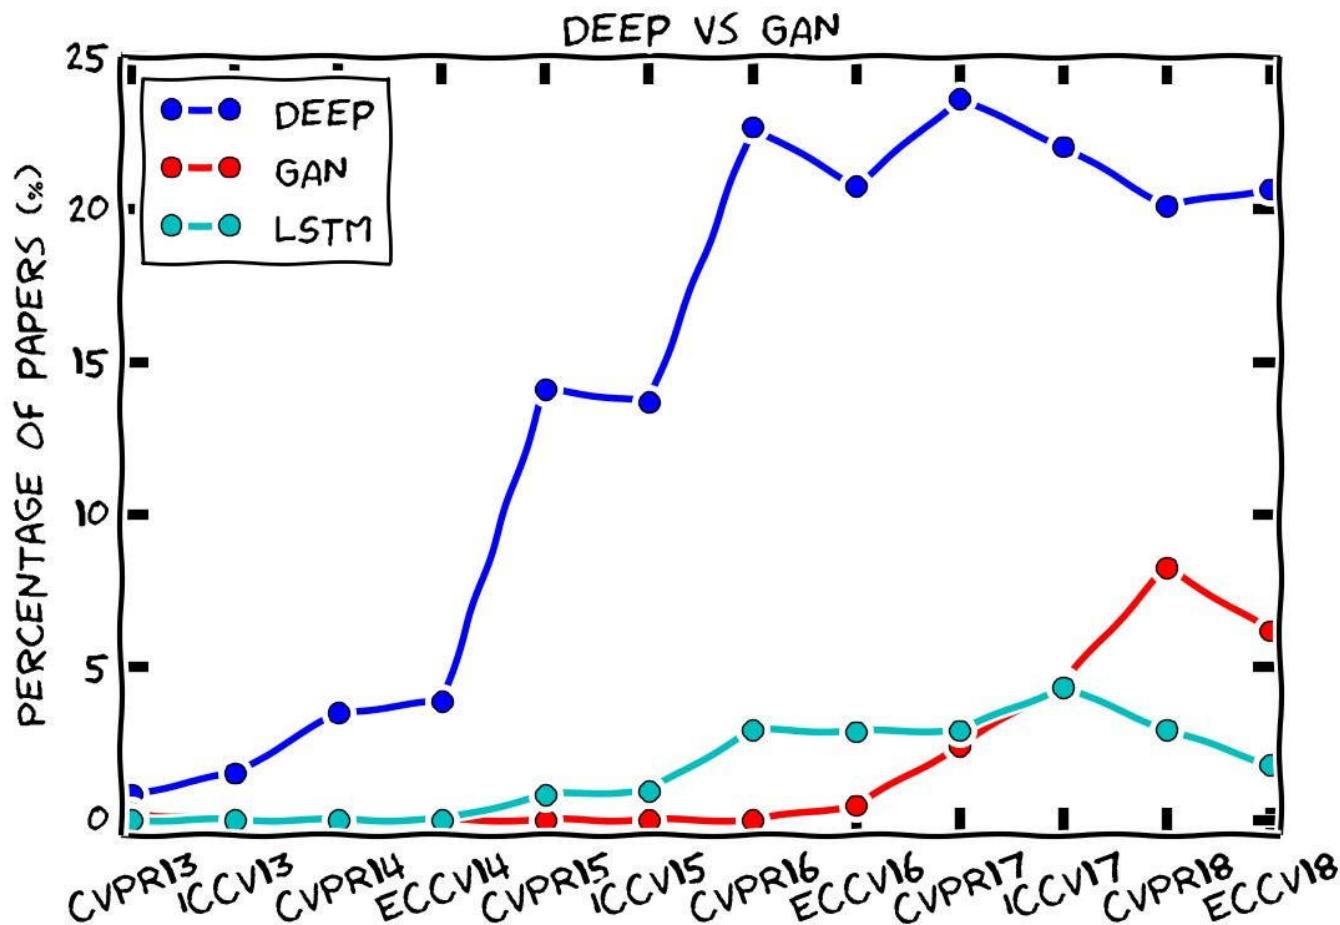
Try to make
Try to make $D(G(z)) = 1$

D(G(z) = 0, D wins, G fails
D(G(z) = 1, D fails, G wins

Generative adversarial network (GAN)



Generative adversarial network (GAN)

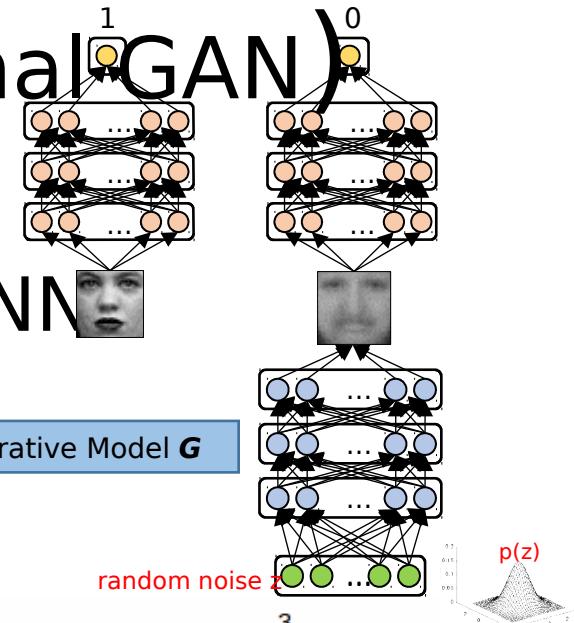


Recap

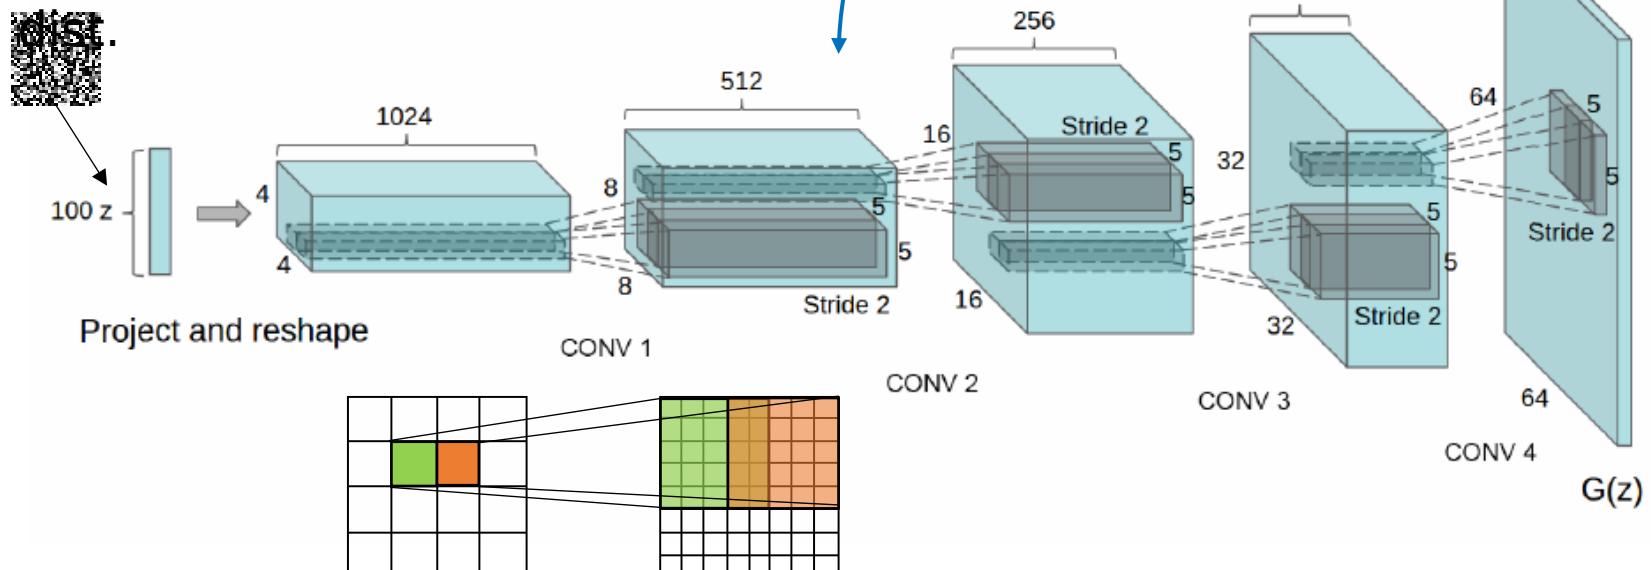
- Supervised learning vs. unsupervised learning
- Generative model
- Autoencoder (sparse AE); generative adversarial network (GAN)

DCGAN (Deep convolutional GAN)

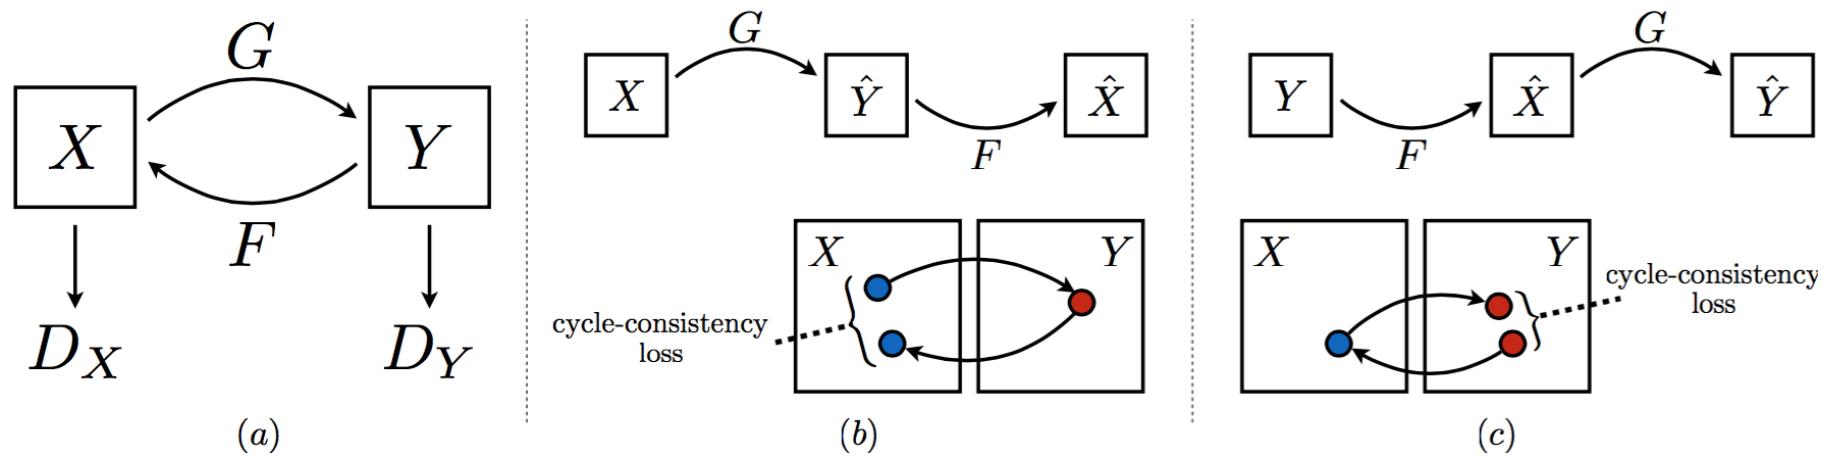
- Replace model's network to CNN
- Example of generator G (same as D)



z : uniform



CycleGAN



CycleGAN

Monet \leftrightarrow Photos



Monet → photo

Zebras \leftrightarrow Horses



zebra → horse

Summer \leftrightarrow Winter



summer → winter



photo → Monet



horse → zebra



winter → summer



Photograph



Monet



Van Gogh



Cezanne



Ukiyo-e

Thank you