

Multi Arm Bandit Problem

April 26, 2019

1 Statistical and Machine Learning (01.113) - HW5 Question 3

In this problem, we will study the multi-armed ($k = 10$) bandit problem.

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
```

```
T = 1000;
```

1.0.1 ϵ -greedy policy

Behave greedily most of the time, but explore once in a while:

$$A_t = \begin{cases} i^* = \arg \max_i Q_t(i) & \text{with probability } 1 - \epsilon \\ j, j \neq i^* & \text{each with probability } \frac{\epsilon}{k-1} \end{cases}$$

Balances exploitation vs exploration, but does not select intelligently between the $k-1$ non-greedy actions.

```
In [2]: # greed and UCB function
def e_greedy(Q,N,t,e):
    """
    Q: 1-d array of action-values
    N: 1-d array of action counts (X)
    t: time (total number of actions taken thus far) (X)
    e: exploration probability

    Returns
    next action to be made
    """
    current_action = np.argmax(Q)
    if np.random.rand()<e: # explore other action with equal probability
        other_actions = [i for i in range(len(Q)) if i != current_action]
        next_action = np.random.choice(other_actions)
    else:
        next_action = current_action
    return next_action
```

1.0.2 Upper confidence bound (UCB)

Select action at time t according to

$$A_t = \arg \max_i \left(Q_t(i) + c \sqrt{\frac{\log t}{N_t(i)}} \right)$$

where $N_t(i)$ denotes the number of times action i has been selected prior to time t , and c is the exploration constant; increasing it favours exploration and decreasing it favours exploitation.

In [3]: *# the denominator means that:*

number of times of A_i increases, we will choose A_i lesser

recall that exploration is choosing other decisions

exploitation is choose current best

```
def UCB(Q,N,t,c):
```

```
    """
```

```
    Q: 1-d array of action-values
```

```
    N: 1-d array of action counts
```

```
    t: time (total number of actions taken thus far)
```

```
    c: exploration constant : increasing it favours exploration \\  
                                   and decreasing it favours exploitation.
```

```
    Returns
```

```
        next action to be made
```

```
    """
```

```
    assert t == sum(N);
```

```
    return np.argmax(Q+c*np.sqrt((np.log(t))/(N)))
```

Write a function that performs one run (1000 time steps), updates Q incrementally and records the reward received at each time step:

$$Q_{n+1} = Q_n + \frac{1}{n+1} (R_{n+1} - Q_n)$$

At each time step, when action a is taken, the reward r is sampled from a normal distribution with mean $true_means[a]$ and standard deviation 1.

In [4]:

```
def test_run(policy,param):
```

```
    true_means=np.random.normal(0,1,10) # true rewards
```

```
    # START, initialize
```

```
    rewards=np.zeros(T+1) # reward tracker
```

```
    Q=np.zeros(10) #action values
```

```
    N = np.zeros(10) # action counts
```

```
    for t in range(T):
```

```
        a=policy(Q,N,t,param) # next action (a)
```

```
        r=np.random.normal(true_means[a],1) # observed reward at (t) when (a)
```

```
        #NewEstimate = OldEstimate + StepSize[Target - OldEstimate]
```

```

        step_size = (1 / (N[a]+1))
        Q[a] += step_size * (r - Q[a])
        N[a] += 1 # update counter
        rewards[t+1]=r # store reward
    return rewards

```

1.0.3 (c) Use the following function to average over 2000 runs and plot the results:

At approximately what value of ϵ does the ϵ -greedy method switch from being better than the greedy policy to being worse than it?

```


In [5]: def simulate_greedy(constants):
    data = []
    for eps in epsilons:
        print("eps: ", eps)
        ave_eg=np.zeros(T+1)
        for i in range(2000):
            eg=test_run(e_greedy, eps) #choose parameter
            ave_eg+=(eg-ave_eg)/(i+1)
        data.append(ave_eg)
    return data

In [10]: time = np.arange(T+1)

    epsilons = [0, 0.1, 0.2, 0.3, 0.4 ,0.5]
    data_greedy = simulate_greedy(epsilons)

    plt.figure(figsize=(10, 6))
    for i in range(len(data_greedy)):
        if i == 0:
            label = "Greedy"
        else:
            label = f"eps: {round(epsilons[i], 2)}"
        plt.plot(time, data_greedy[i], label=label)
        plt.text(time[-1] + 80, data_greedy[i][-1], label, fontsize=12)
    plt.xlabel("Time (t)")
    plt.ylabel("Rewards")
    plt.legend()
    plt.show()

```



We notice that at $\epsilon = 0.3$ has approx. the same performance as greedy, $\epsilon = 0.4$ performs worst, $\epsilon = 0.2$ performs better

```
In [28]: epsilons = [0, 0.27, 0.275, 0.28, 0.285, 0.29, 0.295, 0.3]
data_greedy_new = simulate_greedy(epsilons)

# set 400 as warm up time required to reach steady state
reward_means = [np.mean(rewards[400:]) for rewards in data_greedy_new]
for i in range(len(reward_means)):
    eps = epsilons[i]
    base_mean = reward_means[0]
    r_mean = reward_means[i]
    perform = "Better" if r_mean > base_mean else \
              "Worst" if r_mean < base_mean else ""
    print(f"eps: {eps} Average Rewards: {reward_means[i]} ({perform})")

eps: 0 Average Rewards: 1.0134459686714616 ()
eps: 0.27 Average Rewards: 1.0695985950382605 (Better)
eps: 0.275 Average Rewards: 1.0648032754828465 (Better)
eps: 0.28 Average Rewards: 1.062069782616915 (Better)
eps: 0.285 Average Rewards: 1.033637678572049 (Better)
eps: 0.29 Average Rewards: 1.0366515661034263 (Better)
eps: 0.295 Average Rewards: 1.020058512818882 (Better)
eps: 0.3 Average Rewards: 0.9980903970582163 (Worst)
```